

OPC/AI

My side project started as a football model. It somehow turned into all of this.

- Reddit r/SideProject
 - 5/10
 - #side project #AI # #
 - tactica. Web Match Intelligence Perfor-
mance Intelligence WebSocket
-

I built Perclick. An easy way to match your messaging to your Ads. Would love some feedback.

- Reddit r/SideProject
 - 4/10
 - #Perclick # #
 - Perclick SaaS 1,000 sub-
reddit Allbirds On
-

AI /Agent

I forked Ninfer 3090 and converted it to run on the CMP170HX - doubled my Qwen3.6-35B from llama.cpp

- Reddit r/LocalLLaMA
 - 8/10
 - #Ninfer #CMP170HX # # # LLM
 - Ninfer RTX 3090 CMP170HX 70 SM sm_80 3090 82
SM sm_86 cooperative kernel split-K 16→8→4→2
Blackwell cudaErrorCooperativeLaunchTooLarge Qwen3.6-35B-
A3B int8 KV cache 262K 26GiB llama.cpp 2
llama-swap
-

16 GB VRAM purgatory discussion thread

- Reddit r/LocalLLaMA
 - 7/10
 - # LLM #VRAM #GGUF
-

llm 0.33

- Simon Willison's Blog
 - 6/10
 - #llm # #
 - llm 0.33 OpenAI Python 3.x HTTP httpx httpx2 llm
embed/embed-multi Python --key/key= key llm
prompt -t/--template llm -m gpt-5.6-luna -o reasoning_effort
high --save lhigh llm -t lhigh -t pelican Responses
API reasoning_summary auto/concise/detailed
-

How I built this

- Ben's Bites
 - 5/10
 - #Agent # #Claude
-

More than just code review

- Simon Willison's Blog
 - 4/10
 - #coding-agents #code-review #agentic-engineering
 - coding agents
-

Your own GGUF

- Reddit r/LocalLLaMA
 - 4/10
 - #GGUF #llama.cpp #
 - GGUF Vulkan ROCm llamacpp GGUF Vulkan llama-
quantize ROCm Gemma 4 31B
-

affaan-m/ECC

- GitHub Trending
 - 4/10
 - #Agent #Claude Code #
 - ECC Claude Code Codex Opencode Cursor agent harness
Skills Instincts Memory Security ecc.tools
GitHub App
-

AI /

I benchmark DFlash 2 (PR build) in llama.cpp on Qwen 3.8 27B against all speculative methods for 3 days. 2.26x on 100 real coding prompts, 4.68x with one n-gram drafter on top. Up to 8x on specific cases.

- Reddit r/LocalLLaMA
 - 9/10
 - #DFlash2 #llama.cpp # # #Qwen
 - llama.cpp PR #27342 Qwen3.8 27B DFlash 2 DFlash 2 100 LiveCodeBench 2.26x
-

Fixed the MTP head on Ornith1.5 35B A3B. +3% TPS -33% wall clock

- Reddit r/LocalLLaMA
 - 8/10
 - #LocalLLaMA #Ornith1.5 #MTP # #TPS
-

Freetokens project is impressive

- Reddit r/LocalLLaMA
 - 8/10
 - # # LLM # #FreeToken #
 - FreeTokens RTX 5080 16GB VRAM + DDR6 64GB + AMD Ryzen 9 9950X3D QWEN3.6-35B-A3B NVFP4 20GB 100 tok/s 1028 token
-

Qwen3.8-27B Q6 is a beast at agentic coding

- Reddit r/LocalLLaMA
 - 7/10
 - #Qwen3.8-27B #agentic coding # LLM #RTX 3090
 - Qwen3.8-27B Q6 agentic coding
-

[AINews] 10% worse, 100x cheaper, 10000x faster: Why Simulation is taking over

- Latent Space
- 5/10
- # # #AI
- AI

GLM and I created a llama.cpp fork optimized for AMD GFX906 (Mi50, Mi60, Radeon VII, GCN HIP) - Machine Learning, LLMs, & AI

- Reddit r/LocalLLaMA
 - 5/10
 - #llama.cpp #AMD #GFX906 # #
 - GLM llama.cpp AMD GFX906 Mi50 Mi60 Radeon VII GCN
HIP GPU fork LLM
-

microsoft/onnxruntime

- GitHub Trending
 - 5/10
 - #ONNX Runtime # #
 - ONNX Runtime PyTorch TensorFlow/Keras scikit-learn LightGBM XGBoost
NVIDIA GPU Transformer PyTorch
-

AI

AI

- - 8/10
 - #AI # # #
 - AI 20% AI 2.7 AI 18% 64 45
AI 20% AI “ AI ” AI
-

A third of web pages published since ChatGPT launched were written by AI, study finds

- TechCrunch AI
 - 5/10
 - #AI #
 - ChatGPT AI AI
-

Windows AI PC

-

- 5/10
 - #Windows #AI PC # #
 - Windows 8GB AI Windows WinUI
3 Chromium/WebView2 AI Agent Windows
AI PC
-