

OPC/AI

I'm an engineer who always sucked at marketing. Two weeks of treating it like an engineering problem: 1.8M reddit views, 5x traffic, #6 paid travel app

- Reddit r/SideProject
 - 6/10
 - #AI # #Reddit #SEO
 - \$2 app 75 Waffle House playground Reddit 4 180
 - Store 6
-

I failed at apps for years. Now I've hit 39k downloads. Here's everything I learned about marketing them

- Reddit r/SideProject
 - 6/10
 - #App #ASO #SEO #UGC #
 - iOS 39k App Store " - " "interior
 - design ai - Miya" ASO SEO Google Trend-
 - s 16 UGC Apple Search Ads
-

AI /Agent

Qwen3.8-Flash on RTX3090 + 64GB RAM (but you only need 12GB VRAM)

- Reddit r/LocalLLaMA
 - 8/10
 - # #llama.cpp #Qwen3.8-Flash # #RTX 3090
-

50% tg increase with offloading "hot" experts to VRAM

- Reddit r/LocalLLaMA
 - 8/10
 - #llama.cpp #MoE #VRAM #
-

use llms to auto annotation your dataset locally

- Reddit r/LocalLLaMA
- 8/10
- #LLM # # #

- llmog YOLO GitHub llama.cpp Google Colab vLLM Kaggle API demo

bilawalsidhu/gods-eye-view

- GitHub Trending
- 8/10
- # # # #AI

calesthio/OpenMontage

- GitHub Trending
- 8/10
- # # #Agent
- OpenMontage AGPLv3 12 GitHub 100+ Trending 700+ AI

abi/screenshot-to-code

- GitHub Trending
- 8/10
- # # #AI

tashfeenahmed/freellmapi

- GitHub Trending
- 8/10
- # #LLM API # #

ChromeDevTools/chrome-devtools-mcp

- GitHub Trending
- 8/10
- #Chrome DevTools #MCP #AI

livekit/agents

- GitHub Trending

AI /

zai-org/GLM-5.3 · Hugging Face

- Reddit r/LocalLLaMA
- 9/10
- #GLM-5.3 # # #
- GLM-5.3 GLM-5.2 Z.ai Code Bench GLM-5.2 50% Terminal
Bench 3.0 Agents' Last Exam SOTA CyberGym SOTA exploitation GLM-5.2

[Release] SOTA GGUFs for Qwen3.8-27B: GSQ-RCO at 2.5 to 3.0 bpw

- Reddit r/LocalLLaMA
- 9/10
- #GGUF #Qwen3.8-27B # #GSQ #RCO
- ISTA-DASLab GSQ Gumbel-Softmax Quantization RCO Riemannian
Constrained Optimization Qwen3.8-27B GGUF 2.50/2.75/3.00
bpw 8.4-10.1 GB llama.cpp Ollama LM Studio 3.00 bpw
AIME25 100.00 GPQA-Diamond 88.89 89.90 LiveCodeBench v6
84.57 85.71 2.75 bpw BF16 75.70 vs 74.34 8.4

I benchmarked 9 open models on spotting fake sources during agentic search (DeepSeek V4, Qwen 3.8, Nemotron 3 Ultra)

- Reddit r/LocalLLaMA
- 9/10
- # #benchmark #agentic search #
- “ ” Agentic EchoNet GLM 5.2 1.000
EAS Qwen3.7 Max DeepSeek V4 Flash 15.8%

AtomicChat/Qwen3.8-Flash-Next-GGUF is Really Good

- Reddit r/LocalLLaMA
- 8/10
- #Qwen3.8-Flash-Next #GGUF # #llama.cpp
- AtomicChat Qwen3.8-Flash-Next-GGUF llama.cpp mmap
PLE 106GB 65GB 55GB prefill 500 t/s oMLX
PLE SSD prefill 180 t/s RAM M4 Max 128GB

AI OpenAI

- - 7/10
 - #OpenAI # # #AI
-

Breeze-TTS-2 initial impressions: genuinely ‘frontier’ TTS

- Reddit r/LocalLLaMA
 - 6/10
 - #TTS # #AI
 - Breeze-TTS-2 TTS 7GB BreezeBlue playground
“ ” TTS
-

AI models flub these intelligence tests. Can you fare any better?

- MIT Technology Review AI
 - 6/10
 - #AI # #
-

ROCm 10.0: A Decade of Open Compute, Built for the Age of Agentic AI

- Reddit r/LocalLLaMA
 - 5/10
 - #ROCm #AMD #llama.cpp
 - AMD ROCm 10.0 Agentic AI 7.14 llama.cpp
10.0 PR #27803
-

ds4 branch with GLM 5.3 Flash support

- Reddit r/LocalLLaMA
 - 5/10
 - #ds4 #GLM 5.3 Flash #
 - ds4 GLM 5.3 Flash M4 Max 128GB prompt
ds4 GLM 5.3 Flash
-

Qwen 3.8 Flash Next ngram look up table offloaded to SSD and streamed in SGLang

- Reddit r/LocalLLaMA
 - 4/10
 - #Qwen #SGLang #SSD offload
 - Qwen 3.8 Flash Next ngram SSD SGLang
garnermccloud Hugging Face NVFP4
-

Tencent/Hy4-preview 770B-A49B weight dropped

- Reddit r/LocalLLaMA
 - 4/10
 - # # #MoE #
 - Hy4-preview 770B 49B MoE Hugging
Face tencent/Hy4-preview
-

AI

AI AI | AG&I

- - 5/10
 - #AI # #
-

Coding

- - 5/10
 - #AI Coding # # #TRAE #
-

OpenAI to start showing ads on ChatGPT's free and Go tiers in India

- TechCrunch AI
 - 4/10
 - #OpenAI # #
-